

Metric-learning methods outperform fine-tuning by a large margin; the three variants tie at the top

