

Explanation of Revisions

Anonymized ACL Rolling Review Submission

Previous submission: ACL Rolling Review, May 2026 cycle. The previous version was *desk rejected* before review; no reviewer comments, scores, or meta-review were issued. Consequently there is no point-by-point review response to provide. This document instead summarizes (i) how the desk-rejection causes were remedied and (ii) the substantive changes made since that submission.

1. Remedies for the desk-rejection causes. The May 2026 version was desk rejected for three format violations, all of which are fixed in the present submission:

- **Bibliography style.** The reference list now uses the official ACL style files (`acl.sty` with the `acl_natbib.bst` author-year bibliography style) from the official `acl-org/acl-style-files` repository; the previously used non-conforming style has been removed.
- **Equations out of margins.** All displayed and inline equations have been rechecked; the compiled manuscript contains zero overfull boxes and no content extending beyond the margins.
- **Placement of Limitations.** The Limitations section now appears immediately after the Conclusion, followed by the Ethics Statement, Acknowledgements, and References, as required.

2. Substantive changes since the previous submission. Beyond the format fixes, the manuscript has been substantially rewritten and is no longer a minor revision:

- The evaluation is now built on **ADIBench**, a new multi-granularity benchmark unifying three public Arabic dialect corpora (NADI 2024 18-way, NADI 2024 5-way coarse, and the amgadhasan 5-city corpus) under a single few-shot protocol, rather than a single corpus.
- The proposed method has been replaced by two mechanisms with separate analyses: **LPA** (Linguistic-Prior Aggregation), a test-time prototype aggregation driven by a dialectological distance matrix, and **SCP** (Supervised Contrastive Pre-training) of the encoder, together with an analysis of why the two do not stack.
- **Seven published few-shot baselines** (ProtoNet, MatchingNet, RelationNet, FineTune, Focal, CB, MAML) are re-implemented and evaluated under the unified protocol, replacing the previous baseline set.
- New experiments include **multi-seed results with mean $\pm$ std**, **Macro-F1 in addition to accuracy**, a **class-balanced control** isolating support-set size from class imbalance, and a **support-size sweep** ($k \in \{1, 3, 5, 10, 20\}$) demonstrating that the reported failure mode of fine-tuning is not a small- k artefact.
- The framing has been revised: the paper now positions the contribution as an analysis of when fine-tuning collapses in few-shot long-tail dialect identification together with two complementary, individually effective remedies.

All changes preserve the anonymity of the submission.